

Epistemic Collapse in Multi-Agent Systems

Helivan Calcifer Computing




Figure 1: Epistemic collapse. All agents start with knowledge of correct answers to all M questions in the environment. As system evolves, agent knowledge is updated only via direct observations from the environment and retrieved via top- k_{ctx} relevance $\langle q, k \rangle e^{-c \cdot \text{age}}$. *Left.* Temporal component of memory effects an agent’s ability to answer a question correctly. *Center.* The proportion of forgotten answers versus stale answers is higher in fast forgetting agents. *Right.* A large environment causes epistemic collapse in both fast forgetting and slow forgetting agents.

1 Motivation

Information is being produced and shared more than ever before. The landscape is not only larger but faster: short-form content, the 24-hour news cycle, and algorithmic social media feeds shorten the shelf life of what is current. For example, the near saturation of public discourse by the Epstein files, abruptly displaced by the conflict in the Middle East and the prospect of war with Iran, is a recent reminder that a story dominates collective attention only until the next one arrives — so the facts in circulation are revised more quickly than most of their consumers re-check them.

The cost of recirculating generated content is understood at the model level, where systems trained on their own outputs degrade [1]. It is far less known at the *system* level, where a system of agents that can observe and repeat information can drift away from a changing world’s current truth even when every individual behaves reasonably. Modeling our information landscape as a multi-agent system in a dynamic environment, our goal is to investigate different strategies for maintaining a robust knowledge system and their sensitivity to various environmental and social parameters. Building on our earlier work on monitoring multi-agent systems (*Control Charts for Multi-agent Systems*, arXiv:2605.11135), we study these questions in a real network of LLM agents and in a mechanistically faithful proxy. Our investigation is organized around a failure mode we call *epistemic collapse*.

2 Setting

We model the information landscape as a multi-agent system evolving in discrete steps. It has three components — an environment, a population of agents, and a communication topology that joins them — each with its own levers; particular instantiations (population size, rates, budgets) are reported with each figure.

Environment. The environment is defined as a set of M

question and answer pairs and is the system’s only source of truth. The environment is dynamic in that the correct answer to a question can change with each time step, with a query q ’s answer changing at rate $\lambda_q \sim \text{LogNormal}(\bar{\lambda}, \sigma_\lambda)$. We study the effect of size of the environment M and the average speed of the environment $\bar{\lambda}$ on epistemic collapse.

Agents. Our system is populated by agents each with a fixed context window of size k_{ctx} . Each agent is grounded in that they are instructed to respond to a question with “I don’t know” if there is not enough evidence in their context to respond confidently. The k_{ctx} most semantically and temporally relevant knowledge, as measured by $\langle q, k \rangle e^{-c \cdot \Delta t_q}$, is put in an agent’s context where Δt_q is the elapsed time since receiving the knowledge. An agent’s knowledge is either *firsthand* — obtained by consulting the environment directly — or *secondhand* — received from a peer, and it retains this provenance as it moves: agents relay answers *verbatim*, so a secondhand copy is an exact copy of some earlier firsthand observation and never drifts by paraphrase. Whether secondhand knowledge counts as evidence enough to ground an answer is a policy choice, and the subject of Sec. 3. We study the effect of the context size k_{ctx} and the speed of memory decay c on system dynamics.

Communication. The agents and the environment together form a fixed communication graph, with the environment as a distinguished node — the graph’s only source. Each agent spends a fixed bandwidth of B actions per step across its incident edges: E on its edge to the environment and $K = B - E$ across its peer edges. The two edge types are not symmetric, and the asymmetry is the system’s central constraint. An exchange on an environment edge (a query) returns the current truth, exact and freshly stamped — the only interaction that *creates* current evidence. An exchange on a peer edge (an ask, targeting a question the asker currently lacks) merely *redistributes* existing copies; the reply (possibly an abstention)

inserts into the asker’s memory as received evidence. Whether received evidence may itself ground answers — whether agents may repeat what they have merely heard — is a policy choice, and is the subject of Sec. 3. The communication levers are the graph’s density, the allocation E of bandwidth to the source, and the sharing policy.

Measurement. On a fixed probe panel, every agent is asked every probe question at regular intervals. Each response is exactly one of *correct*, *stale* (an answer, but out of date), or an *abstention*, so $P(\text{correct}) + P(\text{stale}) + P(\text{IDK}) = 1$. We track these at two levels: *agent level* (the average member) and *system level* (whether *no* member can answer, and whether *some* member is correct). The levels nest — system $\text{IDK} \leq \text{agent IDK}$, agent correct $\leq$ system correct — so both can share an axis.

A measured proxy. Our real systems are networks of LLM agents (gpt-4o-mini); real runs cost money, so most experiments run on a proxy with *zero fitted parameters*. Under verbatim relay the LLM’s causal role reduces to a gate — given the retrieved context, emit the stored payload or abstain — and that gate is measurable: probing the model across its (finitely many, up to exchangeability and similarity bucketing) context states yields a near-deterministic step function, with essentially no leakage of the model’s parametric knowledge. Replacing the LLM call with this measured table, in the same code path, reproduces the real system’s statistics across two full factor grids with median deviation 0.01 on all four metrics (Fig. 2). Results below run on this validated proxy except where marked real.

2.1 Epistemic collapse

We say an agent has *epistemically collapsed* on a question set when it can no longer ground answers and produces only abstentions; the *system* has collapsed when no member can answer. Collapse has a quieter twin: an agent whose evidence has expired but still surfaces it answers *confidently and wrongly* — availability survives, accuracy dies. Both matter, and every mechanism below prices them differently.

Figure 1 shows collapse from the best possible start: every agent warm-started with current, firsthand knowledge of everything. Accuracy falls from ≈ 1 to a plateau set by memory decay (~ 0.61 at $c = 0.05$, ~ 0.14 at $c = 0.5$; left). Conditioned on being wrong, the fast-decay system fails purely by *forgetting* (100% abstention); the slow-decay system’s errors are $\sim 90\%$

forgetting and $\sim 10\%$ *confident staleness* — rising to $\sim 25\%$ in the early transient — with the stale share concentrated on the fastest-revising questions. Abstention is insensitive to how fast any fact revises, but highly sensitive to how many facts there are: steady-state $P(\text{IDK})$ climbs from 0.18 to 0.90 as the environment grows from 25 to 400 questions at fixed budget (right) — breadth, not speed, silences the population, because every new fact’s insertions crowd the others out of retrievable memory. The same regimes appear in the real gpt-4o-mini factor grids (project artifacts): decay ($c : 0.05 \rightarrow 1.0$ takes agent $\text{IDK } 0.35 \rightarrow 0.89$), source starvation (environment bandwidth $E : 6 \rightarrow 1$ takes system $\text{IDK } 0 \rightarrow 0.29$), environment speed ($\bar{\lambda} : 0.001 \rightarrow 0.1$ takes system accuracy $0.93 \rightarrow 0.20$ without moving abstention).

The rest of this note studies the two natural *mitigations*, both forms of communication: (1) **sharing secondhand information** (mimesis) — letting agents repeat answers they received rather than observed; and (2) **books** — a shared artifact agents write to and read from. Both run on the validated platform of Sec. 2.

3 Mitigation 1: sharing secondhand information

The mimesis grid repeats the factor design of Sec. 2 with one change: received answers may ground an answer (still verbatim — agents repeat, they do not paraphrase). Figure 3 overlays the two grids at the agent level.

Mimesis raises the floor. At baseline, agent-level IDK drops $0.68 \rightarrow 0.26$ and agent accuracy rises $0.23 \rightarrow 0.48$. The communication factor switches from a null lever to a live one: agent accuracy is 0.38 at degree 2 vs. 0.56 on the mesh, where without mimesis the same sweep is flat.

But not the ceiling. System-level accuracy does not improve — 0.72 (no mimesis) vs. 0.66 (mimesis) at baseline — because the system level asks whether *anyone* holds current evidence, and repetition redistributes evidence without creating it. Mimesis pays at the system level only where redistribution is the binding constraint: in the observation-starved cell ($E = 1$) system accuracy rises $0.46 \rightarrow 0.64$ and system IDK falls $0.29 \rightarrow 0.02$. The regime it cannot rescue is fast forgetting ($c = 1.0$: system $\text{IDK } 0.33$ vs. 0.31) — copies crowd out as fast as originals.




Figure 2: Validation of the proxy: the simulation matches the real system, per factor, with zero fitted parameters. Each panel: that factor’s three settings $\times$ four metrics, plotted as simulation (horizontal, median with 10-seed IQR) against real (vertical, median with seed range). The chain is: measured answering gate $P(\text{response category} \mid \text{context state})$ + verbatim relay + mechanical retrieval. Median $|\Delta| = 0.01$ in every factor; the same holds for the mimesis grid (not shown).

4 Mitigation 2: books

Mimesis is peer-to-peer. The second mitigation is a shared artifact: a *book* any agent can read for one action, fed only by deliberate, costly writes — an agent attempts a write with probability p , an attempt consumes its entire step, and a write commits after W attempts, merging the writer’s current grounded knowledge (newest evidence wins). Reads ground answers like observations. We sweep the toy platform (10 seeds per point) in a scarce-observation regime where books have room to matter.

The book’s value is a race between its write rate and the world’s revision rate. The natural unit is the expected write rate $\omega = p/W$ (completed writes per agent per step). The book’s accuracy gain over the no-books floor is large when $\bar{\lambda} \ll \omega$ ($+0.48$ at $\omega = 10/100$ steps, $\bar{\lambda} = 0.001$) and gone when $\bar{\lambda}/\omega \gtrsim 1$: every condition lands on the floor by $\bar{\lambda} = 0.3$ (Fig. 4, left). Availability behaves completely differently: $P(\text{IDK})$ is set by ω ($0.43/0.63/0.83$ for $\omega = 10, 1, 0.2$ per 100 steps, vs. 0.93 without books) and is *flat* in $\bar{\lambda}$ (Fig. 4, right). Staleness absorbs the entire loss: the book keeps everyone answering while the answers quietly expire — a fast world does not silence the library, it turns it into a confident source of expired facts.

ω is the sufficient statistic. Holding p fixed and varying W (constant budget tax) reproduces the ordering, confirming W is a pure freshness dial with diminishing returns once the book turns over faster than the world. Conversely, holding ω fixed while varying p by $4\text{--}10\times$ moves the curves by at most 0.06 on either metric, against between- ω gaps up to 0.22 . The residual is monotone in p with the sign a budget tax predicts (high- p writers lose more steps to attempts and deposit slightly staler content, $\approx 0.04\text{--}0.06$ accuracy at slow speeds). For design purposes: quote a book policy by its expected writes; the attempt probability is a second-order tax.

Takeaway and next steps. Communication — secondhand sharing or shared artifacts — reliably buys *availability*, and buys *accuracy* only while the underlying evidence is younger than the world’s revision clock: mimesis is capped by what the

population has observed, and a book by how often anyone can afford to write. As the environment speeds up, the failure mode is not silence but confident staleness, which availability metrics cannot see. Next: (1) “dated editions” — books exposing their write time so readers can abstain on stale entries, trading availability back for accuracy; (2) strategy-sensitivity sweeps on the validated platform; (3) the monitoring question: does a control chart on the probe panel see confident staleness coming, given it presents as drift, not shift?

References

- [1] I. Shumailov, Z. Shumaylov, Y. Zhao, N. Papernot, R. Anderson, Y. Gal. AI models collapse when trained on recursively generated data. *Nature* 631, 755–759 (2024).




Figure 3: Sharing secondhand information raises the floor, not the ceiling (agent level; same factor grids and protocol as Sec. 2, mimesis vs. no mimesis overlaid). Mimesis roughly halves abstention and doubles accuracy for the average member at baseline and brings the communication graph into play; system-level accuracy (not shown) stays environment-capped: 0.72 vs. 0.66 at baseline.




Figure 4: A book buys availability at any speed; its accuracy dies at $\bar{\lambda}/\omega \sim 1$. Toy platform, no mimesis, scarce observation; conditions labeled by expected write rate $\omega = p/W$. Left: accuracy with the book converges to the no-books floor as $\bar{\lambda}$ rises, for every ω . Right: $P(\text{IDK})$ is set by ω and essentially flat in $\bar{\lambda}$.